

## CST4070 - Applied Data Analytics - Tools, Practical Big Data Handling, Cloud Distribution Summative assessment – Component 2 Individual report - Dragomir Nedev M00724882

##Problem definition Three datasets are available: bike\_journeys, bike\_stations and LondonCensus. Spatial granularity: each bike station. Temporal granularity: one hour time slot. Goal: predicting the total number of bikes rented in each bike station with the temporal granularity of one hour time slot.

### ##Preprocecssing Importing the datasets

```
library(data.table)
```

```
## Warning: package 'data.table' was built under R version 3.6.2
```

```
bike_journeys = fread('bike_journeys.csv')
bike_stations = fread('bike_stations.csv')
census = fread('London_census.csv')
```

### Exploring the datasets

```
head(bike_journeys)
```

```
##      Journey_Duration Journey_ID End_Date End_Month End_Year End_Hour End_Minute
## 1:           2040           953      19       9       17       18         0
## 2:           1800          12581      19       9       17       15        21
## 3:           1140           1159      15       9       17       17         1
## 4:            420           2375      14       9       17       12        16
## 5:           1200          14659      13       9       17       19        33
## 6:           1320           2351      14       9       17       14        53
##      End_Station_ID Start_Date Start_Month Start_Year Start_Hour Start_Minute
## 1:           478      19       9       17       17         26
## 2:           122      19       9       17       14         51
## 3:           639      15       9       17       16         42
## 4:           755      14       9       17       12          9
## 5:           605      13       9       17       19        13
## 6:           514      14       9       17       14        31
##      Start_Station_ID
## 1:           251
## 2:           550
## 3:           212
## 4:           163
## 5:            36
## 6:           589
```

```
head(bike_stations)
```

```
##      Station_ID Capacity Latitude Longitude      Station_Name
## 1:           1      19 51.52916 -0.109970      River Street , Clerkenwell
## 2:           2      37 51.49961 -0.197574      Phillimore Gardens, Kensington
## 3:           3      32 51.52128 -0.084605      Christopher Street, Liverpool Street
## 4:           4      23 51.53006 -0.120973      St. Chad's Street, King's Cross
## 5:           5      27 51.49313 -0.156876      Sedding Street, Sloane Square
## 6:           6      18 51.51812 -0.144228      Broadcasting House, Marylebone
```

```
head(census)
```

```
##      WardCode      WardName      borough NESW AreaSqKm      lon
## 1: E05000026      Abbey Barking and Dagenham East      1.3 0.077935
## 2: E05000027      Alibon Barking and Dagenham East      1.4 0.148270
## 3: E05000028      Becontree Barking and Dagenham East      1.3 0.118957
## 4: E05000029      Chadwell Heath Barking and Dagenham East      3.4 0.139985
## 5: E05000030      Eastbrook Barking and Dagenham East      3.5 0.173581
## 6: E05000031      Eastbury Barking and Dagenham East      1.4 0.105683
##      lat IncomeScor LivingEnSc NoEmployee GrenSpace PopDen BornUK NotBornUK
## 1: 51.53971      0.27      42.76      7900      19.6 9884.6      5459      7327
## 2: 51.54559      0.28      27.96      800      22.4 7464.3      7824      2561
## 3: 51.55453      0.25      31.59      1100      3.0 8923.1      8075      3470
## 4: 51.58475      0.27      34.78      1700      56.4 2970.6      7539      2482
## 5: 51.55365      0.19      21.25      4000      51.1 3014.3      8514      1992
## 6: 51.53590      0.27      31.16      1000      18.1 8357.1      7880      3744
##      NoCTFtoH NoDwelling NoFlats NoHouses NoOwndDwel MedHPrice
## 1:      0.1      4733      3153      1600      1545      177000
## 2:      0.1      4045      574      3471      1849      160000
## 3:      0.1      4378      837      3541      2093      170000
## 4:      0.4      4050      1400      2662      2148      195000
## 5:      0.5      3976      742      3235      2646      191750
## 6:      0.0      4321      933      3388      1913      167250
```

Importing libraries which will check and plot heatmap of missing values

```
library(Rcpp)
```


```
## Warning: package 'Rcpp' was built under R version 3.6.2
```

```
library(Amelia)
```

```
## Warning: package 'Amelia' was built under R version 3.6.2
```

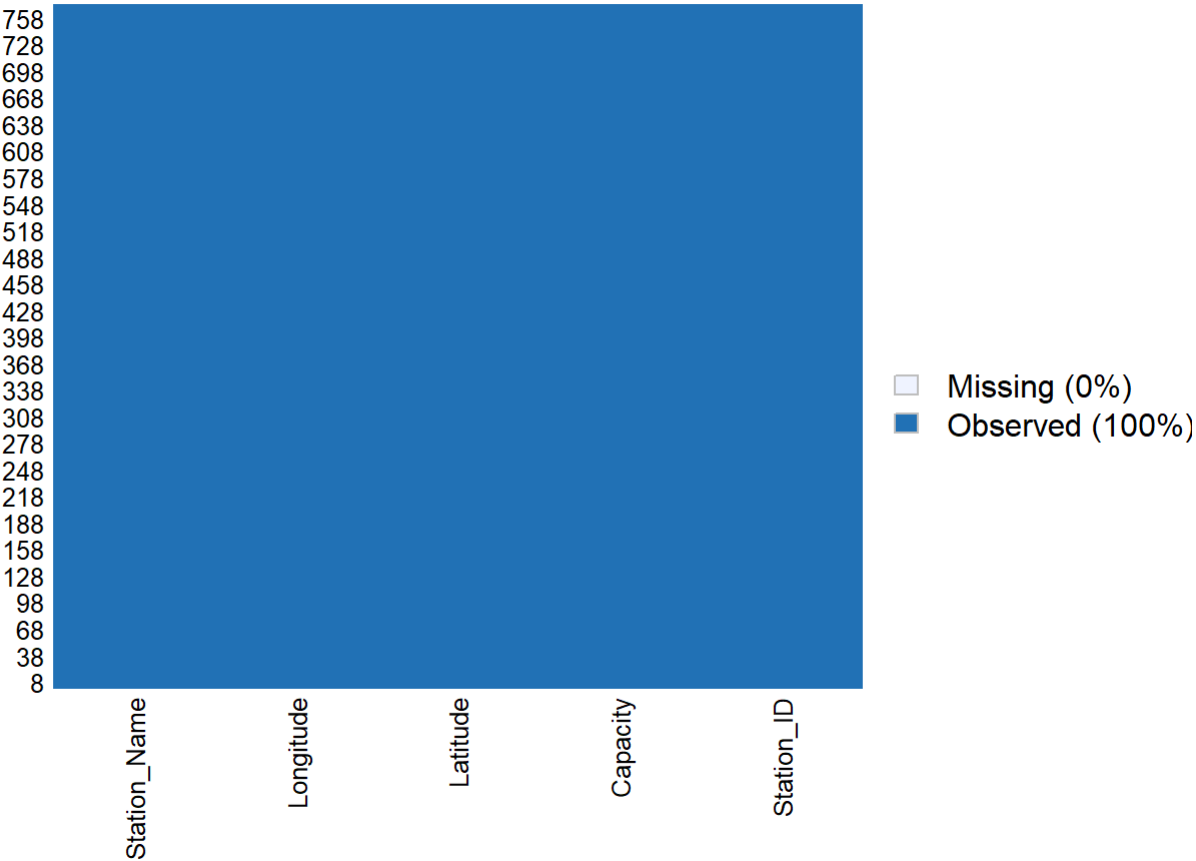
```
## ##
## ## Amelia II: Multiple Imputation
## ## (Version 1.7.6, built: 2019-11-24)
## ## Copyright (C) 2005-2020 James Honaker, Gary King and Matthew Blackwell
## ## Refer to http://gking.harvard.edu/amelia/ for more information
## ##
```

```
missmap(bike_journeys)
```




```
missmap(bike_stations)
```

### Missingness Map



```
missmap(census)
```



missing data, in the datasets, means there would not be NaN values.

Checking consistency between bike\_journeys and bike\_stations. We have to join this datasets based on Start\_Station\_ID and StationID so we need to check whether they contain the same values.

```
length(unique(bike_journeys$Start_Station_ID))

## [1] 779

length(unique(bike_journeys$End_Station_ID))

## [1] 779

length(unique(bike_stations$Station_ID))

## [1] 773

length(unique(intersect(bike_stations$Station_ID, bike_journeys$Start_Station_ID)))

## [1] 771
```

Bike\_journeys dataset contains 779 unique stations (the same number for end stations and start stations). Bike\_stations dataset contains 773 unique stations. Both datasets have 771 matching unique stations which means that we will exclude data for 8 stations.

## ##Hypotheses

H1. Bikes demand is higher during peak hours. H2. Bikes demand have a daily trend. H3. Higher demand of bikes rented at stations which are close to central London. H4. Higher demand of bikes rented where is high employment rate. H5. Higher demand of bikes rented where is high population density. H6. Higher demand of bikes rented where is high percentage of green space. H7. Higher demand of bikes rented in deprived areas. H8. Higher demand of bikes rented in poor areas. H9. Higher demand of bikes rented where is high immigration rate. H10. Higher demand of bikes rented where is high flats rate. H11. Higher demand of bikes rented where is low number of owned properties rate.

## ##Metrics

- bike\_rides. Number of rides would be our dependant variable that we need to predict
- Start\_hour. Indicate the hour when the journey started. Linked to H1.
- Start\_Day. Indicate the day when the journey started. Linked to H2.
- finalRatioEmployee. Ratio of people who are employed. NoEmployee over PopDen times AreaSqKm. Linked to H4.
- PopDen. Population divided by the ward area. Linked to H5.
- GrenSpace. Percentage of green space associated with the ward. Linked to H6.
- LivingEnSc. Quality of the local environment. The more deprived is an area, the higher the score. Linked to H7.
- IncomeScor. Proportion of the population experiencing deprivation relating to low income. Higher score means lower income and poorer areas. Linked to H8.
- MedHPrice. Median house price. The lower median means the poorer areas. Linked to H8.
- RatioCTFtoH. Ratio of properties in council tax band F-H (the highest median house price). The lower score means the poorer areas. Linked H8.
- RatioBornUK. Ratio of people who were born in the UK. It is defined as NotBornUK over BornUK plus NotBornUK. Linked to H9.
- FlatsRate. Ratio of flats. It is defined as NoFlats over NoHouses. Linked to H10.
- RatioOwndDwel. Ratio of owned properties in each ward. It is defined as NoOwndDwel over NoDwelling. Linked to H11.

## ##Data processing

Due to the fact that the census data holds the record of longitude and latitude of the ward and the bike\_station dataset, contains the coordinates of the bike stations, we need to calculate the nearest distance. Importing library "geosphere" will help us calculate the distance between the locations from the two datasets

```
library(geosphere)
```

```
## Warning: package 'geosphere' was built under R version 3.6.2
```

```
distance <- distm(bike_stations[, 4:3], census[, 6:7])
```

```
distance_calc <- cbind(bike_stations, census[apply(distance, 1, which.min),])
View(distance_calc)
```

Renaming the column Start\_Station\_ID to match Station\_ID, so we could merge the data

```
colnames(bike_journeys)[colnames(bike_journeys) == "Start_Station_ID"] <- "Station_ID"
```

After we are done the the transformations of the location we can merge the datasets

```
total <- merge(bike_journeys,distance_calc,by = "Station_ID")
```

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 3.6.2
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:data.table':  
##  
##      between, first, last
```

```
## The following objects are masked from 'package:stats':  
##  
##      filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##      intersect, setdiff, setequal, union
```

Combining the different data fields into one

```
total$Journey_date <- as.Date(with(total, paste(Start_Year, Start_Month, Start_Date ,sep="-")),  
"%y-%m-%d")
```

```
total2 <- total %>% group_by(Start_Hour, Station_ID, Journey_date) %>% summarise(bike_rides = n  
( ))
```

```
View(total2)
```

```
total2 <- left_join(total, total2, by=c("Station_ID", "Start_Hour", "Journey_date")) %>% rowwise  
( )
```

The data frame needs to be transformed into a datatable, before extracting the final dataset

```
setDT(total2)
```

The data needs to be transformed from the format:

<Journey\_Duration, Journey\_ID, End\_Date, End\_Month, End\_Year, End\_Hour, End\_Minute, End\_Station\_ID, Start\_Date, Start\_Month, Start\_Year, Start\_Hour, Start\_Minute, Start\_Station\_ID> <Station\_ID, Capacity, Latitude, Longitude, Station\_Name> <WardCode, WardName, Borough, NESW, AreaSqKm, Ion, lat, IncomeScor, LivingEnSc, NoEmployee, GrenSpace, PopDen, BornUK, NotBornUK, NoCTFtoH, NoDwelling, NoFlats, NoHouses, NoWndDwel, MedHPrice>

Into the format:

<bike\_rides, Station\_ID, Start\_Date, Start\_Hour, MedHPrice, finalRatioEmployee, IncomeScor, LivingEnSc, GrenSpace, RatioBornUK, RatioCTFtoH, RatioOwndDwel, FlatsRate>

```
final = total2[, .(bike_rides, Station_ID, Start_Date,
  Start_Hour, MedHPrice,
  finalRatioEmployee=NoEmployee/(PopDen*AreaSqKm), IncomeScor, LivingEnSc,
  GrenSpace, RatioBornUK=BornUK/(BornUK+NotBornUK), RatioCTFtoH=NoCTFtoH/(NoDwe
lling),
  RatioOwndDwel=NoOwndDwel/NoDwelling, FlatsRate=NoFlats/(NoFlats+NoHouses))]
```

str(final)

```
## Classes 'data.table' and 'data.frame': 1530240 obs. of 13 variables:
## $ bike_rides : int 2 4 1 1 4 7 8 10 1 3 ...
## $ Station_ID : int 1 1 1 1 1 1 1 1 1 1 ...
## $ Start_Date : int 17 14 18 13 19 15 15 19 17 13 ...
## $ Start_Hour : int 12 7 6 6 6 9 8 8 19 19 ...
## $ MedHPrice : int 455000 455000 455000 455000 455000 455000 455000 455000 455000 45
5000 ...
## $ finalRatioEmployee: num 3.82 3.82 3.82 3.82 3.82 ...
## $ IncomeScor : num 0.21 0.21 0.21 0.21 0.21 0.21 0.21 0.21 0.21 0.21 ...
## $ LivingEnSc : num 51 51 51 51 51 ...
## $ GrenSpace : num 9.3 9.3 9.3 9.3 9.3 9.3 9.3 9.3 9.3 9.3 ...
## $ RatioBornUK : num 0.615 0.615 0.615 0.615 0.615 ...
## $ RatioCTFtoH : num 0.00424 0.00424 0.00424 0.00424 0.00424 ...
## $ RatioOwndDwel : num 0.276 0.276 0.276 0.276 0.276 ...
## $ FlatsRate : num 0.905 0.905 0.905 0.905 0.905 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

Summerising the information from the final dataset

```
summary(final)
```



```
##      bike_rides      Station_ID      Start_Date      Start_Hour
## Min.   : 1.000    Min.   : 1.0    Min.   : 1.0    Min.   : 0.00
## 1st Qu.: 3.000    1st Qu.:163.0    1st Qu.: 7.0    1st Qu.: 9.00
## Median : 5.000    Median :333.0    Median :13.0    Median :14.00
## Mean   : 8.576    Mean   :366.8    Mean   :13.8    Mean   :13.76
## 3rd Qu.: 9.000    3rd Qu.:570.0    3rd Qu.:19.0    3rd Qu.:18.00
## Max.   :182.000    Max.   :826.0    Max.   :31.0    Max.   :23.00
##      MedHPrice      finalRatioEmployee      IncomeScor      LivingEnSc
## Min.   : 188000    Min.   : 0.1321    Min.   :0.0100    Min.   :22.05
## 1st Qu.: 362500    1st Qu.: 0.5446    1st Qu.:0.0900    1st Qu.:43.29
## Median : 455000    Median : 1.4114    Median :0.1700    Median :48.34
## Mean   : 559274    Mean   : 5.4905    Mean   :0.1766    Mean   :48.36
## 3rd Qu.: 652500    3rd Qu.: 3.8660    3rd Qu.:0.2400    3rd Qu.:53.64
## Max.   :1750000    Max.   :50.5540    Max.   :0.4400    Max.   :68.06
##      GrenSpace      RatioBornUK      RatioCTFtoH      RatioOwndDwel
## Min.   : 0.00    Min.   :0.3543    Min.   :2.661e-05    Min.   :0.1380
## 1st Qu.: 7.50    1st Qu.:0.4785    1st Qu.:2.491e-03    1st Qu.:0.2167
## Median :13.50    Median :0.5521    Median :4.613e-03    Median :0.2707
## Mean   :17.61    Mean   :0.5333    Mean   :5.683e-03    Mean   :0.2803
## 3rd Qu.:25.00    3rd Qu.:0.5955    3rd Qu.:8.481e-03    3rd Qu.:0.3352
## Max.   :69.10    Max.   :0.7112    Max.   :1.794e-02    Max.   :0.5476
##      FlatsRate
## Min.   :0.5423
## 1st Qu.:0.8397
## Median :0.8928
## Mean   :0.8733
## 3rd Qu.:0.9480
## Max.   :0.9794
```

In a few of the variables it could be seen that they are not normally distributed, which indicates that they have to be transformed in to log value.

```
final$bike_rides = log10(final$bike_rides + min(final[bike_rides!=0]$bike_rides))
final$RatioBornUK = log10(final$RatioBornUK + min(final[RatioBornUK!=0]$RatioBornUK))
final$RatioCTFtoH = log10(final$RatioCTFtoH + min(final[RatioCTFtoH!=0]$RatioCTFtoH))
```

Standardising the data

```
mydata_std = as.data.table(scale(final) )
summary(mydata_std)
```

```
##      bike_rides      Station_ID      Start_Date      Start_Hour
##  Min.      :-1.46429    Min.      :-1.5385    Min.      :-1.51647    Min.      :-2.80363
##  1st Qu.: -0.59281    1st Qu.: -0.8571    1st Qu.: -0.80589    1st Qu.: -0.96936
##  Median : -0.08303    Median : -0.1421    Median : -0.09531    Median :  0.04969
##  Mean   :  0.00000    Mean   :  0.0000    Mean   :  0.00000    Mean   :  0.00000
##  3rd Qu.:  0.55922    3rd Qu.:  0.8548    3rd Qu.:  0.61528    3rd Qu.:  0.86492
##  Max.    :  4.21399    Max.    :  1.9315    Max.    :  2.03645    Max.    :  1.88396
##      MedHPrice      finalRatioEmployee      IncomeScor      LivingEnSc
##  Min.      :-1.1803    Min.      :-0.4830    Min.      :-1.64916    Min.      :-3.24909
##  1st Qu.: -0.6255    1st Qu.: -0.4458    1st Qu.: -0.85726    1st Qu.: -0.62599
##  Median : -0.3315    Median : -0.3677    Median : -0.06537    Median : -0.00232
##  Mean   :  0.0000    Mean   :  0.0000    Mean   :  0.00000    Mean   :  0.00000
##  3rd Qu.:  0.2964    3rd Qu.: -0.1464    3rd Qu.:  0.62754    3rd Qu.:  0.65222
##  Max.    :  3.7852    Max.    :  4.0623    Max.    :  2.60727    Max.    :  2.43307
##      GrenSpace      RatioBornUK      RatioCTFtoH      RatioOwndDwel
##  Min.      :-1.2211    Min.      :-2.2525    Min.      :-3.9194    Min.      :-1.9282
##  1st Qu.: -0.7011    1st Qu.: -0.6027    1st Qu.: -0.3728    1st Qu.: -0.8612
##  Median : -0.2852    Median :  0.2625    Median :  0.1893    Median : -0.1297
##  Mean   :  0.0000    Mean   :  0.0000    Mean   :  0.0000    Mean   :  0.0000
##  3rd Qu.:  0.5121    3rd Qu.:  0.7397    3rd Qu.:  0.7469    3rd Qu.:  0.7441
##  Max.    :  3.5694    Max.    :  1.9145    Max.    :  1.4342    Max.    :  3.6237
##      FlatsRate
##  Min.      :-3.5828
##  1st Qu.: -0.3632
##  Median :  0.2117
##  Mean   :  0.0000
##  3rd Qu.:  0.8088
##  Max.    :  1.1487
```

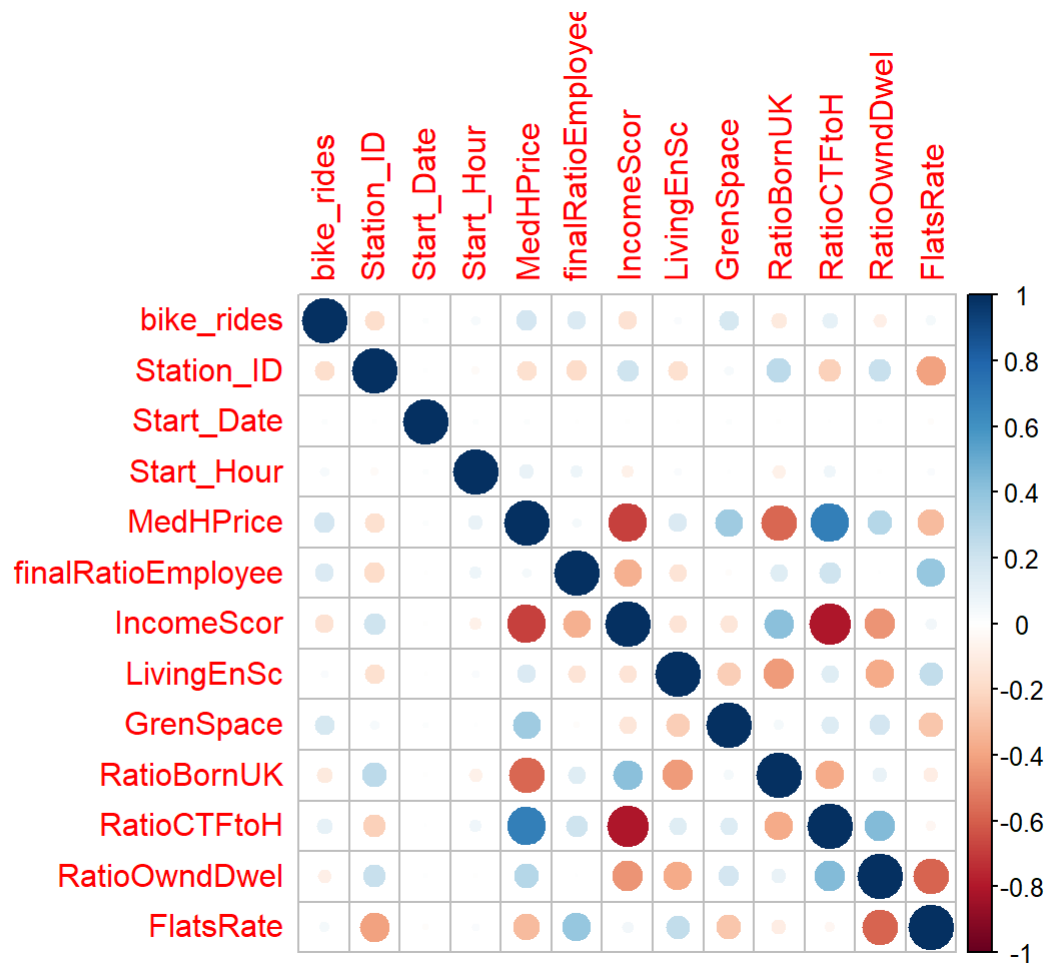
Checking for multicollinearity.

```
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 3.6.2
```

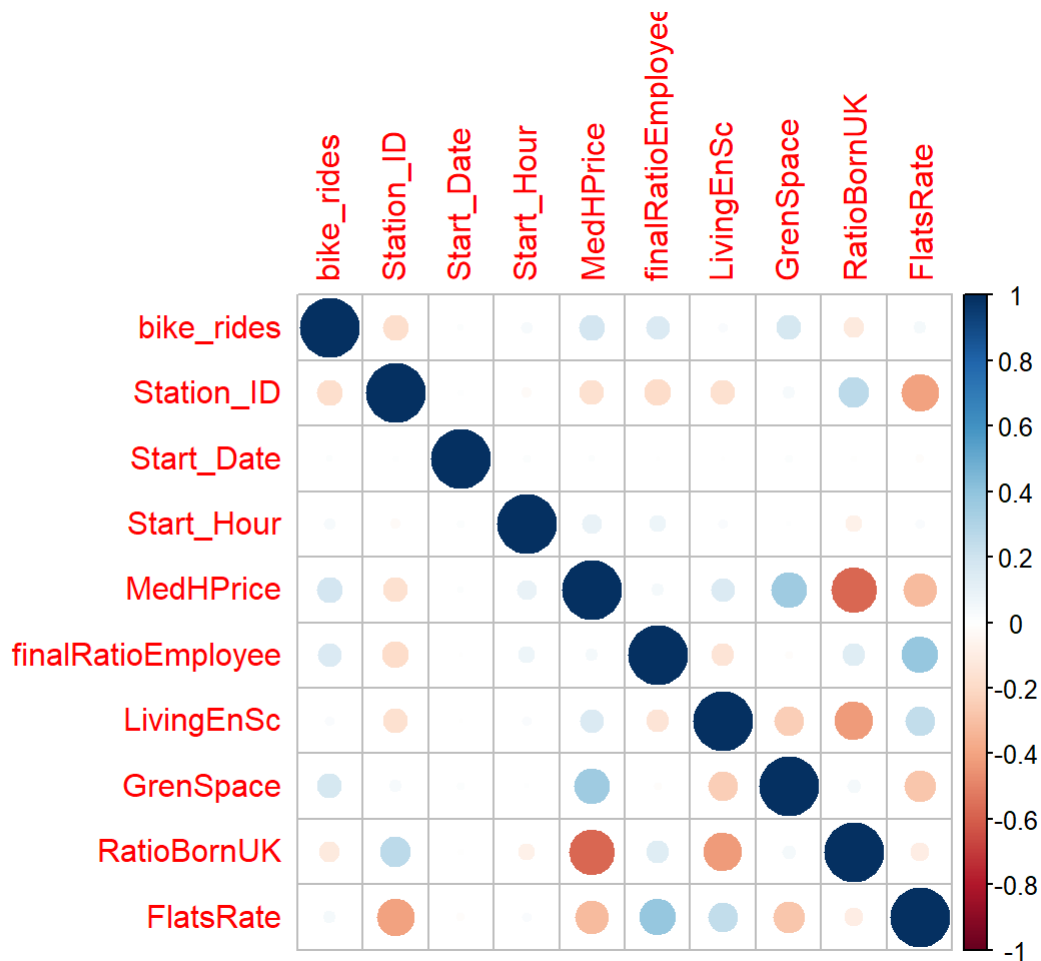
```
## corrplot 0.84 loaded
```

```
corrplot(cor(mydata_std))
```



There is high correlation between RatioCTFtoH, RatioOwndDwel and IncomeScor so they will be removed from the model. Again checking multicollinearity.

```
mydata_std$RatioCTFtoH = NULL
mydata_std$RatioOwndDwel = NULL
mydata_std$IncomeScor = NULL
corrplot(cor(mydata_std))
```



##Algorithms Linear regression model needs to implemented as part of the final goal

```
set.seed(0)
trainIdx = sample(1:nrow(mydata_std), 0.75*nrow(mydata_std))
train = mydata_std[trainIdx]
test = mydata_std[-trainIdx]
lr = lm(bike_rides ~ ., data=train)
train_preds = predict(lr, train)
test_preds = predict(lr, test)
```

Printing the R2 scores

```
print(paste("R2 on train:", cor(train_preds, train$bike_rides)^2))
```

```
## [1] "R2 on train: 0.0925138096168877"
```

```
print(paste("R2 on test:", cor(test_preds, test$bike_rides)^2))
```

```
## [1] "R2 on test: 0.0933658930141483"
```

##Data understanding Plotting the beta coefficients ot understand the model.


```
lr = lm(bike_rides ~ ., data=mydata_std)
summary(lr)
```

```
##
## Call:
## lm(formula = bike_rides ~ ., data = mydata_std)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4997 -0.6818 -0.0653  0.5922  4.1429
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.825e-13  7.700e-04   0.00      1
## Station_ID     -1.316e-01  8.970e-04 -146.68 <2e-16 ***
## Start_Date      9.129e-03  7.703e-04   11.85 <2e-16 ***
## Start_Hour      1.142e-02  7.768e-04   14.70 <2e-16 ***
## MedHPrice       3.821e-02  1.328e-03   28.77 <2e-16 ***
## finalRatioEmployee 1.579e-01  9.661e-04  163.42 <2e-16 ***
## LivingEnSc       3.590e-02  9.199e-04   39.02 <2e-16 ***
## GrenSpace       1.694e-01  9.063e-04  186.91 <2e-16 ***
## RatioBornUK     -7.413e-02  1.185e-03  -62.54 <2e-16 ***
## FlatsRate       -2.711e-02  1.170e-03  -23.18 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9525 on 1530230 degrees of freedom
## Multiple R-squared:  0.09273,    Adjusted R-squared:  0.09272
## F-statistic: 1.738e+04 on 9 and 1530230 DF,  p-value: < 2.2e-16
```

```
library(ggplot2)
```


```
## Warning: package 'ggplot2' was built under R version 3.6.2
```

```
ggplot(, aes(x = names(lr$coefficients), y=lr$coefficients)) +
  geom_bar(stat="identity") +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust = 0.5)) +
  xlab("coefficient") +
  ylab("normalised value")
```



Checking the multicollinearity of the data


```
corrplot(cor(mydata_std))
```



Plotting histograms to prove H1 and H2

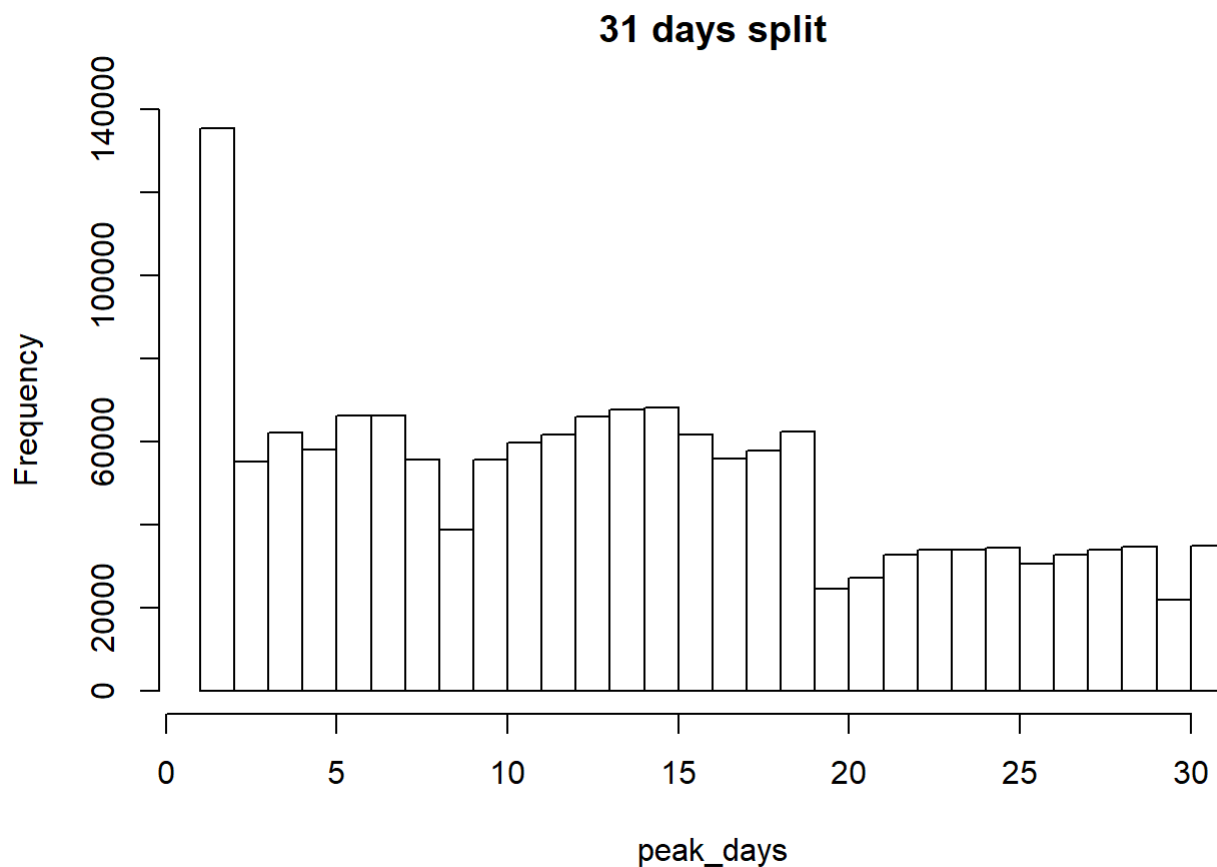
```
peak_hours <- total12$Start_Hour
hist(peak_hours, breaks = 24, main = "24 hours split")
```

## 24 hours split



```
peak_days <- total2$Start_Date  
hist(peak_days, breaks = 31, main = "31 days split")
```





##Main findings H1. Bikes demand is higher durin peak hours. TRUE. As seen from vis 24 hour split, where we can see there are clear high deman during peak hours (07-09:00 and 16-18:00), which prooves our hypothesis.

H2. Bikes demand have a daily trend. TRUE. As seen in vis 31 days split, there is higher demand in the first half of the month, than the second half.

H3. Higher demand of bikes rented at stations which are close to central London. Cannot be falsified due to the the fact that the data needs to be standardised and the values for the locations is not numeric

H4. Higher demand of bikes rented where is high employment rate. TRUE.

H5. Higher demand of bikes rented where is high population density. Cannot be falsified due to multicollinearity.

H6. Higher demand of bikes rented where is high percentage of green space. TRUE. We can see that the bike\_rides are fairly high correlated to the zones with high concentration of green spaces.

H7. Higher demand of bikes rented in deprived areas. FALSE

H8. Higher demand of bikes rented in poor areas. TRUE. Lower demand in wealthier zones.

H9. Higher demand of bikes rented where is high immigration rate. TRUE. We can see that the bike\_rides are fairly high correlated to the zones where there are people who are predominantly born in UK

H10. Higher demand of bikes rented where is high flats rate. FALSE.

H11. Higher demand of bikes rented where is low number of owned properties rate. Cannot be falsified due to multicollinearity of OwndDwelRate

##Limitaions

1. The short period of time, reviewed in the dataset, does not allow us to do perfected model. More months would give us better predictions
2. Multicollinearity of some of the features reduces the accuracy of the model
3. Introducing weather data would further improve our model as we would be able to take external factors.